Real Time Data Streaming and Decision Making

Opex Analytics

Table of Contents

Introduction

The aim of this project is to develop and prototype a real time data streaming and decision making solution. As a working example, we will use the bike sharing data. The first step towards achieving this is to develop and prototype a solution that reads real-time streaming data, stores it and presents it in a format ready for consumption by analytical processes. This is then followed by using the data streamed to build analytical solutions to make real time decisions. For this working example, the questions answered would be:

- Where are there are bike stock outs?
- Where will there be stockouts in the next x minutes.
- What is the probability there will be a bike available for me at station x in y minutes?
- Where should I send my trucks with extra bikes to restock?

Data

Data Sources

The following is a table comparing the various vendors:

Vendor	Locations	Refresh Interval(s)	Historic Data (since)
Bike Town	Portland, OR	60	N.A.
Capital Bike Share	Washington, DC	10	2010 Q4
Citi Bike	New York, NY	10	2013 Q3
CoGo Bikes	Columbus, OH	30	N.A.
Divy Bikes	Chicago, IL	30	2013 Q3
Nice Ride	Minnesota, MN	10	2010 June
The Hubway	Boston, MA	10	2011-2013

All the above mentioned vendors publish data in General Bikeshare Feed Specification (GBFS) format. All files shared in the GBFS format are JSON files with utf-8 encoding. Every JSON file presented in this specification contains the same common header information at the top level of the JSON response object:

Field Name	Required	Defines
last_updated	Yes	Integer POSIX timestamp indicating the last time the data in this feed was updated.
ttl	Yes	Integer representing the number of seconds before the data is refreshed (0 if the data should always be refreshed).
data	Yes	JSON hash containing the data fields for this response.

Files Shared by Vendors

This specification defines the following files along with their associated content:

File Name	Required	Defines
gbfs.json	Optional	Auto-discovery file that links to all of the other files published by the system. This file is optional, but highly recommended.
system_information.json	Yes	Describes the system including System operator, System location, year implemented, URLs, contact info, time zone
station_information.json	Yes	Mostly static list of all stations, their capacities and locations
station_status.json	Yes	Number of available bikes and docks at each station and station availability
free_bike_status.json	Optional	Describes bikes that are available in non station-based systems
system_hours.json	Optional	Describes the hours of operation for the system
system_calendar.json	Optional	Describes the days of operation for the system
system_regions.json	Optional	Describes the regions the system is broken up into
system_pricing_plans.json	Optional	Describes the system pricing
system_alerts.json	Optional	Describes current system alerts

gbfs.json

The following fields are attributes in the gbfs.json file:

Field Name	Required	Defines
language	Yes	The language that all of the contained files will be published in. This language must match the value in the system_information file.
-feeds	Yes	An array of all the feeds that are published by this auto-discovery file.
-name	Yes	Key identifying the type of feed this is (e.g. "system_information", "station_information").
-url	Yes	Full URL for the feed.

Example:


```
{
  "last_updated": 1500042893,
  "ttl": 10,
  "data": {
    "en": {
      "feeds": [{
        "name": "system_information",
        "url": "https://gbfs.citibikenyc.com/gbfs/en/system_information.json"
      }, {
        "name": "station_status",
        "url": "https://gbfs.citibikenyc.com/gbfs/en/station_status.json"
      }, {
        "name": "system_alerts",
        "url": "https://gbfs.citibikenyc.com/gbfs/en/system_alerts.json"
      }, {
        "name": "system_regions",
        "url": "https://gbfs.citibikenyc.com/gbfs/en/system_regions.json"
      }, {
        "name": "station_information",
        "url": "https://gbfs.citibikenyc.com/gbfs/en/station_information.json"
      }]
    }
  }
}
```

system_information.json

The following fields are all attributes within the main "data" object for this feed.

Field Name	Required	Defines
system_id	Yes	ID field - identifier for this bike share system. This should be globally unique. In addition, this value is intended to remain the same over the life of the system
language	Yes	An IETF language tag indicating the language that will be used throughout the rest of the files. This is a string that defines a single language tag only.
name	Yes	Full name of the system to be displayed to customers
short_name	Optional	Optional abbreviation for a system
operator	Optional	Name of the operator of the system
url	Optional	The URL of the bike share system. The value is a fully qualified URL that includes http:// or https://, and any special characters in the URL must be correctly escaped.
purchase_url	Optional	A fully qualified URL where a customer can purchase a membership or learn more about purchasing memberships
start_date	Optional	String in the form YYYY-MM-DD representing the date that the system began operations
phone_number	Optional	A single voice telephone number for the specified system. This field is a string value that presents the telephone number as typical for the system's service area.

email	Optional	A single contact email address for customers to address questions about the system
timezone	Yes	The time zone where the system is located.
license_url	Optional	A fully qualified URL of a page that defines the license terms for the GBFS data for this system, as well as any other license terms the system would like to define (including the use of corporate trademarks, etc)

Example:


```
{
  "last_updated": 1500046512,
  "ttl": 10,
  "data": {
    "system_id": "NYC",
    "language": "en",
    "name": "Citi Bike",
    "short_name": "Citi Bike",
    "operator": "Motivate International, Inc.",
    "url": "http://www.citibikenyc.com",
    "purchase_url": "http://www.citibikenyc.com/",
    "start_date": "2013-05-01",
    "phone_number": "1-855-BIKE-311",
    "email": "customerservice@motivateco.com",
    "license_url": "",
    "timezone": "America/New_York"
  }
}
```

station_information.json

All stations contained in this list are considered public (ie, can be shown on a map for public use). If there are private stations (such as Capital Bikeshare's White House station) these should not be exposed here and their status should not be included in station_status.json.

Field Name	Required	Defines
stations	Yes	Array that contains one object per station in the system as defined below
- station_id	Yes	Unique identifier of a station.
- name	Yes	Public name of the station
- short_name	No	Short name or other type of identifier, as used by the data publisher
- lat	Yes	The latitude of station.
- lon	Yes	The longitude of station.
- address	Optional	Valid street number and name where station is located.
- cross_street	Optional	Cross street of where the station is located.
- region_id	Optional	ID of the region where station is located.
- post_code	Optional	Postal code where station is located
- rental_methods	Optional	Array of enumerables containing the payment methods accepted at this station.
- capacity	Optional	Number of total docking points installed at this station, both available and unavailable

Example:


```
{
  "last_updated": 1500046939,
  "ttl": 10,
  "data": {
    "stations": [{
      "station_id": "72",
      "name": "W 52 St & 11 Ave",
      "short_name": "6926.01",
      "lat": 40.76727216,
      "lon": -73.99392888,
      "region_id": 71,
      "rental_methods": ["CREDITCARD", "KEY"],
      "capacity": 39,
      "eightd_has_key_dispenser": false
    }, {
      "station_id": "79",
      "name": "Franklin St & W Broadway",
      "short_name": "5430.08",
      "lat": 40.71911552,
      "lon": -74.00666661,
      "region_id": 71,
      "rental_methods": ["CREDITCARD", "KEY"],
      "capacity": 33,
      "eightd_has_key_dispenser": false
    }]
  }
}
```

station_status.json

Field Name	Required	Defines
stations	Yes	Array that contains one object per station in the system as defined below.
- station_id	Yes	Unique identifier of a station.
- num_bikes_available	Yes	Number of bikes available for rental.
- num_bikes_disabled	Optional	Number of disabled bikes at the station.
- num_docks_available	Yes	Number of docks accepting bike returns.
- num_docks_disabled	Optional	Number of empty but disabled dock points at the station. This value remains as part of the spec as it is possibly useful during development
- is_installed	Yes	1/0 boolean - is the station currently on the street.
- is_renting	Yes	1/0 boolean - is the station currently renting bikes.
- is_returning	Yes	1/0 boolean - is the station accepting bike.
- last_reported	Yes	Integer POSIX timestamp indicating the last time this station reported its status to the backend

Example:


```
{
  "last_updated": 1500047771,
  "ttl": 10,
  "data": {
    "stations": [{
      "station_id": "72",
      "num_bikes_available": 9,
      "num_bikes_disabled": 3,
      "num_docks_available": 27,
      "num_docks_disabled": 0,
      "is_installed": 1,
      "is_renting": 1,
      "is_returning": 1,
      "last_reported": 1500047719,
      "eightd_has_available_keys": false
    }, {
      "station_id": "79",
      "num_bikes_available": 0,
      "num_bikes_disabled": 1,
      "num_docks_available": 32,
      "num_docks_disabled": 0,
      "is_installed": 1,
      "is_renting": 1,
      "is_returning": 1,
      "last_reported": 1500047540,
      "eightd_has_available_keys": false
    }
```

Technologies

Data Ingestion

Technology	Pros	Cons
Kafka	Highly available, redundant, scalable, one-to-many messaging.	Relatively new, lack of commercial support, no built in connectors to Hadoop products.
Flume	Stable, well-established, natively supported by Hadoop.	One-to-one messaging, not redundant.

Data Processing

Technology	Pros	Cons
Storm	Very low latency, well-established.	No guarantee that data will be processed only once.
Spark-streaming	Data is processed reliably, interfaces seamlessly with Spark.	Micro-batching instead of true streams, moderate latency.
Flink	Very low latency, data is processed reliably.	No commercial support.

Data Storage

Technology	Pros	Cons
Cassandra	Availability, easy maintenance, good with large scale.	No multiple secondary indexes, no dynamic querying on columns, slow on-the fly aggregations.
MongoDB	Flexible, highly available, good for real time systems.	No foreign key.
Hbase	Optimized for read, strict consistency, fast r/w with scalability.	Only range based scan, not good for data aggregation.

References

- https://github.com/BetaNYC/Bike-Share-Data-Best-Practices/wiki/Bike-Share-Data-Systems
- https://www.motivateco.com/use-our-data/
- https://resources.zaloni.com/blog/top-streaming-technologies-for-data-lakes-and-real-time-data